

# 1. Student Information

- **Name:** Apoorv Gupta
- **PRN:** 21070126018
- **Batch:** AIML - A1
- **Git Repo:** [GitHub Repository \(https://github.com/erApoorvGupta/NLP\\_assignments\)](https://github.com/erApoorvGupta/NLP_assignments).

## 1.1. Introduction

## 1.2. Importing the libraries

```
In [ ]: # Importing the libraries
        # Preprocessing the data using NLTK

        # Importing the libraries/////
        import nltk
        from nltk.tokenize import word_tokenize
        from nltk.stem import WordNetLemmatizer
        import pandas as pd
        nltk.download('all')
```

## 1.3. Importing the dataset

```
In [2]: df = pd.read_csv(r"/kaggle/input/amazon-fine-food-reviews/Reviews.csv")
df.head()
```

Out[2]:

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	Summary	Text
0	1	B001E4KFG0	A3SGXH7AUHU8GW	delmartian	1	1	5	1303862400	Good Quality Dog Food	I have bought several of the Vitality canned d...
1	2	B00813GRG4	A1D87F6ZCVE5NK	dll pa	0	0	1	1346976000	Not as Advertised	Product arrived labeled as Jumbo Salted Peanut...
2	3	B000LQOCH0	ABXLMWJIXXAIN	Natalia Corres "Natalia Corres"	1	1	4	1219017600	"Delight" says it all	This is a confection that has been around a fe...
3	4	B000UA0QIQ	A395BORC6FGVXV	Karl	3	3	2	1307923200	Cough Medicine	If you are looking for the secret ingredient i...
4	5	B006K2ZZ7K	A1UQRSCLF8GW1T	Michael D. Bigham "M. Wassir"	0	0	5	1350777600	Great taffy	Great taffy at a great price. There was a wid...

## 2. Data Preprocessing

```
In [4]: df.dropna(inplace=True)
```

```
In [5]: df = df[['Text', 'Score']].dropna()
```

```
In [6]: df.drop_duplicates(subset=['Text', 'Score'], keep='first', inplace=True)
```

```
In [29]: def mark_sentiment(Score):
    if(Score<=3):
        return 0
    else:
        return 1
```

```
In [ ]: df['sentiment']=df['Score'].apply(mark_sentiment)
```

```
In [10]: df.drop(['Score'],axis=1,inplace=True) #
```

```
In [13]: df['sentiment'].value_counts()
```

```
Out[13]: 1    306812  
        0     86849  
        Name: sentiment, dtype: int64
```

```
In [14]: df = df.iloc[:40000]
```

## 2.1. Lemmatization and Tokenization

```
In [15]: # Initialize the Lemmatizer  
lemmatizer = WordNetLemmatizer()  
  
# Defining a function to tokenize and Lemmatize the text  
  
def tokenize_and_lemmatize(text):  
    tokens = word_tokenize(text)  
    lemmatized_tokens = [lemmatizer.lemmatize(token) for token in tokens]  
    return " ".join(lemmatized_tokens)
```

```
In [ ]: import nltk  
nltk.download('stopwords')  
nltk.download('wordnet')  
! unzip /usr/share/nltk_data/corpora/wordnet.zip -d /usr/share/nltk_data/corpora/
```

```
In [17]: # Applying the function to the text column  
df['Text'] = df['Text'].apply(tokenize_and_lemmatize)
```

## 3. Data Cleaning

### 3.1 Remove stopwords, Remove symbols, Remove URLs

```
In [18]: # Data Cleansing: Remove stopwords, remove symbols, remove URLs
```

```
# Importing the libraries
import re
from nltk.corpus import stopwords

stop_words = set(stopwords.words('english'))
```

```
In [19]: # Defining a function to clean the text
```

```
def clean_Text(text):
    # Remove URLs
    text = re.sub(r'http\S+', '', text)
    # Remove symbols and numbers
    text = re.sub(r'^\w\s', '', text)
    # Remove stopwords
    text = " ".join([word for word in text.split() if word.lower() not in stop_words])

    # Remove excess whitespaces
    text = ' '.join(text.split())

    # Replace abbreviations (you can add more if needed)
    text = re.sub(r"won't", "will not", text)
    text = re.sub(r"can't", "cannot", text)

    # Fix contractions
    text = re.sub(r"n't", " not", text)
    text = re.sub(r"'re", " are", text)
    text = re.sub(r"'s", " is", text)
    text = re.sub(r"'d", " would", text)
    text = re.sub(r"'ll", " will", text)
    text = re.sub(r"'t", " not", text)
    text = re.sub(r"'ve", " have", text)
    return text
```

```
In [20]: df.rename(columns={'Text': 'text'}, inplace=True)
```

```
In [21]: # Applying the clean Text function to the Text column
df['text'] = df['text'].apply(clean_Text)

# Displaying the first 5 rows of the dataset
df.head()
```

Out[21]:

	text	sentiment
0	bought several Vitality canned dog food produc...	1
1	Product arrived labeled Jumbo Salted Peanuts p...	0
2	confection ha around century light pillowy cit...	1
3	looking secret ingredient Robitussin believe f...	0
4	Great taffy great price wa wide assortment yum...	1

```
In [22]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 100000 entries, 0 to 114536
Data columns (total 2 columns):
#   Column      Non-Null Count  Dtype
---  -
0   text        100000 non-null object
1   sentiment   100000 non-null int64
dtypes: int64(1), object(1)
memory usage: 2.3+ MB
```

```
In [23]: import tensorflow as tf
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Embedding, LSTM, Dense

# Define parameters
batch_size_1 = 4
max_sequence_length_1 = 50
embedding_dim_1 = 50
max_words_1 = 10000
lstm_units_1 = 32

# Tokenize the text
tokenizer = Tokenizer(num_words=max_words_1)
tokenizer.fit_on_texts(df['text'])
sequences = tokenizer.texts_to_sequences(df['text'])
x = pad_sequences(sequences, maxlen=max_sequence_length_1)
y = df['sentiment']

x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state=42)

# Define the first model (1st set of results)
model_1 = Sequential()
model_1.add(Embedding(max_words_1, embedding_dim_1, input_length=max_sequence_length_1))
model_1.add(LSTM(lstm_units_1))
model_1.add(Dense(1, activation='sigmoid'))
model_1.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])

# Train the first model
model_1.fit(x_train, y_train, batch_size=batch_size_1, epochs=20)

# Evaluate the first model
y_pred_1 = model_1.predict(x_test)
y_pred_1 = (y_pred_1 > 0.5) # Threshold for binary classification

# Generate a classification report for the first model
report_1 = classification_report(y_test, y_pred_1)
```

```
/opt/conda/lib/python3.10/site-packages/tensorflow_io/python/ops/__init__.py:98: UserWarning: unable to load libtensorflow_io_plugins.so: unable to open file: libtensorflow_io_plugins.so, from paths: ['/opt/conda/lib/python3.10/site-packages/tensorflow_io/python/ops/libtensorflow_io_plugins.so']
caused by: ['/opt/conda/lib/python3.10/site-packages/tensorflow_io/python/ops/libtensorflow_io_plugins.so: undefined symbol: _ZN3tsl6StatusC1EN10tensorflow5error4CodeESt17basic_string_viewIcSt11char_traitsIcEENS_14SourceLocationE']
  warnings.warn(f"unable to load libtensorflow_io_plugins.so: {e}")
/opt/conda/lib/python3.10/site-packages/tensorflow_io/python/ops/__init__.py:104: UserWarning: file system plugins are not loaded: unable to open file: libtensorflow_io.so, from paths: ['/opt/conda/lib/python3.10/site-packages/tensorflow_io/python/ops/libtensorflow_io.so']
caused by: ['/opt/conda/lib/python3.10/site-packages/tensorflow_io/python/ops/libtensorflow_io.so: undefined symbol: _ZTVN10tensorflow13GcsFileSystemE']
  warnings.warn(f"file system plugins are not loaded: {e}")
```

```
Epoch 1/20
20000/20000 [=====] - 147s 7ms/step - loss: 0.3252 - accuracy: 0.8611
Epoch 2/20
20000/20000 [=====] - 113s 6ms/step - loss: 0.2588 - accuracy: 0.8920
Epoch 3/20
20000/20000 [=====] - 113s 6ms/step - loss: 0.2200 - accuracy: 0.9102
Epoch 4/20
20000/20000 [=====] - 111s 6ms/step - loss: 0.1858 - accuracy: 0.9259
Epoch 5/20
20000/20000 [=====] - 113s 6ms/step - loss: 0.1532 - accuracy: 0.9415
Epoch 6/20
20000/20000 [=====] - 112s 6ms/step - loss: 0.1221 - accuracy: 0.9541
Epoch 7/20
20000/20000 [=====] - 108s 5ms/step - loss: 0.0931 - accuracy: 0.9659
Epoch 8/20
20000/20000 [=====] - 110s 6ms/step - loss: 0.0681 - accuracy: 0.9753
Epoch 9/20
20000/20000 [=====] - 109s 5ms/step - loss: 0.0511 - accuracy: 0.9816
Epoch 10/20
20000/20000 [=====] - 113s 6ms/step - loss: 0.0372 - accuracy: 0.9872
Epoch 11/20
20000/20000 [=====] - 114s 6ms/step - loss: 0.0286 - accuracy: 0.9903
Epoch 12/20
20000/20000 [=====] - 110s 5ms/step - loss: 0.0236 - accuracy: 0.9926
Epoch 13/20
20000/20000 [=====] - 110s 5ms/step - loss: 0.0198 - accuracy: 0.9935
Epoch 14/20
20000/20000 [=====] - 112s 6ms/step - loss: 0.0165 - accuracy: 0.9948
Epoch 15/20
20000/20000 [=====] - 111s 6ms/step - loss: 0.0146 - accuracy: 0.9956
Epoch 16/20
20000/20000 [=====] - 109s 5ms/step - loss: 0.0139 - accuracy: 0.9954
Epoch 17/20
20000/20000 [=====] - 111s 6ms/step - loss: 0.0113 - accuracy: 0.9962
Epoch 18/20
20000/20000 [=====] - 110s 6ms/step - loss: 0.0109 - accuracy: 0.9967
Epoch 19/20
20000/20000 [=====] - 108s 5ms/step - loss: 0.0117 - accuracy: 0.9966
Epoch 20/20
625/625 [=====] - 2s 2ms/step - loss: 0.0096 - accuracy: 0.9973
```



```
In [24]: y_pred_1 = model_1.predict(x_test)
y_pred_1 = (y_pred_1 > 0.5) # Threshold for binary classification

# Generate a classification report for the first model
report_1 = classification_report(y_test, y_pred_1)
print("Classification Report for Model 1:")
print(report_1)
```

625/625 [=====] - 2s 2ms/step

Classification Report for Model 1:

	precision	recall	f1-score	support
0	0.68	0.65	0.66	4597
1	0.90	0.91	0.90	15403
accuracy			0.85	20000
macro avg	0.79	0.78	0.78	20000
weighted avg	0.85	0.85	0.85	20000

```
In [25]: # Define parameters
batch_size_2 = 8
max_sequence_length_2 = 30
embedding_dim_2 = 30
max_words_2 = 25000
lstm_units_2 = 32

# Tokenize the text
tokenizer = Tokenizer(num_words=max_words_2)
tokenizer.fit_on_texts(df['text'])
sequences = tokenizer.texts_to_sequences(df['text'])
x = pad_sequences(sequences, maxlen=max_sequence_length_2)
y = df['sentiment']

x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state=42)

# # Define the model
# model_2 = Sequential()
# model_2.add(Embedding(max_words_2, embedding_dim_2, input_length=max_sequence_length_2))
# model_2.add(LSTM(lstm_units_2, return_sequences=True))
# model_2.add(LSTM(lstm_units_2))
# model_2.add(Dense(1, activation='sigmoid'))

# # Compile the model
# model_2.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])

# # Train the model (assuming you have 'sentiment' as your target column)
# model_2.fit(x, y, batch_size=batch_size_2, epochs=5)
model_2 = Sequential()
model_2.add(Embedding(max_words_2, embedding_dim_2, input_length=max_sequence_length_2))
model_2.add(LSTM(lstm_units_2, return_sequences=True))
model_2.add(LSTM(lstm_units_2))
model_2.add(Dense(1, activation='sigmoid'))
model_2.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])

# Train the second model
model_2.fit(x_train, y_train, batch_size=batch_size_2, epochs=20)

# Evaluate the second model
y_pred_2 = model_2.predict(x_test)
y_pred_2 = (y_pred_2 > 0.5) # Threshold for binary classification

# Generate a classification report for the second model
report_2 = classification_report(y_test, y_pred_2)
```

```
Epoch 1/20
10000/10000 [=====] - 94s 9ms/step - loss: 0.3421 - accuracy: 0.8529
Epoch 2/20
10000/10000 [=====] - 70s 7ms/step - loss: 0.2692 - accuracy: 0.8880
Epoch 3/20
10000/10000 [=====] - 71s 7ms/step - loss: 0.2207 - accuracy: 0.9101
Epoch 4/20
10000/10000 [=====] - 70s 7ms/step - loss: 0.1762 - accuracy: 0.9309
Epoch 5/20
10000/10000 [=====] - 71s 7ms/step - loss: 0.1362 - accuracy: 0.9475
Epoch 6/20
10000/10000 [=====] - 70s 7ms/step - loss: 0.0999 - accuracy: 0.9632
Epoch 7/20
10000/10000 [=====] - 70s 7ms/step - loss: 0.0705 - accuracy: 0.9748
Epoch 8/20
10000/10000 [=====] - 70s 7ms/step - loss: 0.0495 - accuracy: 0.9827
Epoch 9/20
10000/10000 [=====] - 70s 7ms/step - loss: 0.0336 - accuracy: 0.9885
Epoch 10/20
10000/10000 [=====] - 71s 7ms/step - loss: 0.0252 - accuracy: 0.9920
Epoch 11/20
10000/10000 [=====] - 72s 7ms/step - loss: 0.0182 - accuracy: 0.9939
Epoch 12/20
10000/10000 [=====] - 72s 7ms/step - loss: 0.0145 - accuracy: 0.9952
Epoch 13/20
10000/10000 [=====] - 71s 7ms/step - loss: 0.0117 - accuracy: 0.9962
Epoch 14/20
10000/10000 [=====] - 71s 7ms/step - loss: 0.0093 - accuracy: 0.9970
Epoch 15/20
10000/10000 [=====] - 71s 7ms/step - loss: 0.0079 - accuracy: 0.9974
Epoch 16/20
10000/10000 [=====] - 70s 7ms/step - loss: 0.0072 - accuracy: 0.9977
Epoch 17/20
10000/10000 [=====] - 70s 7ms/step - loss: 0.0053 - accuracy: 0.9983
Epoch 18/20
10000/10000 [=====] - 72s 7ms/step - loss: 0.0049 - accuracy: 0.9984
Epoch 19/20
10000/10000 [=====] - 75s 7ms/step - loss: 0.0052 - accuracy: 0.9983
Epoch 20/20
10000/10000 [=====] - 71s 7ms/step - loss: 0.0046 - accuracy: 0.9986
625/625 [=====] - 3s 3ms/step
```

```
In [26]: print("Classification Report for Model 2:")  
print(report_2)
```

Classification Report for Model 2:

	precision	recall	f1-score	support
0	0.62	0.61	0.61	4597
1	0.88	0.89	0.89	15403
accuracy			0.82	20000
macro avg	0.75	0.75	0.75	20000
weighted avg	0.82	0.82	0.82	20000



```

In [27]: import matplotlib.pyplot as plt
from sklearn.metrics import confusion_matrix, roc_auc_score, roc_curve, auc

# Function to plot the confusion matrix
def plot_confusion_matrix(y_true, y_pred, title):
    cm = confusion_matrix(y_true, y_pred)
    plt.figure(figsize=(6, 4))
    plt.imshow(cm, interpolation='nearest', cmap=plt.cm.Blues)
    plt.title(title)
    plt.colorbar()
    classes = ["Negative", "Positive"] # Assuming 0 is negative and 1 is positive
    tick_marks = [0, 1]
    plt.xticks(tick_marks, classes)
    plt.yticks(tick_marks, classes)
    plt.xlabel('Predicted label')
    plt.ylabel('True label')
    for i in range(2):
        for j in range(2):
            plt.text(j, i, format(cm[i, j], 'd'), horizontalalignment="center", color="white" if cm[i, j] > cm.max() / 2 else "black")
    plt.show()

# Function to plot ROC AUC curve
def plot_roc_auc(y_true, y_score, title):
    fpr, tpr, thresholds = roc_curve(y_true, y_score)
    roc_auc = auc(fpr, tpr)
    plt.figure(figsize=(6, 4))
    plt.plot(fpr, tpr, color='darkorange', lw=2, label='ROC curve (area = {:.2f})'.format(roc_auc))
    plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
    plt.xlim([0.0, 1.0])
    plt.ylim([0.0, 1.05])
    plt.xlabel('False Positive Rate')
    plt.ylabel('True Positive Rate')
    plt.title(title)
    plt.legend(loc='lower right')
    plt.show()


# Assuming you have already trained model_1 and model_2 as mentioned earlier

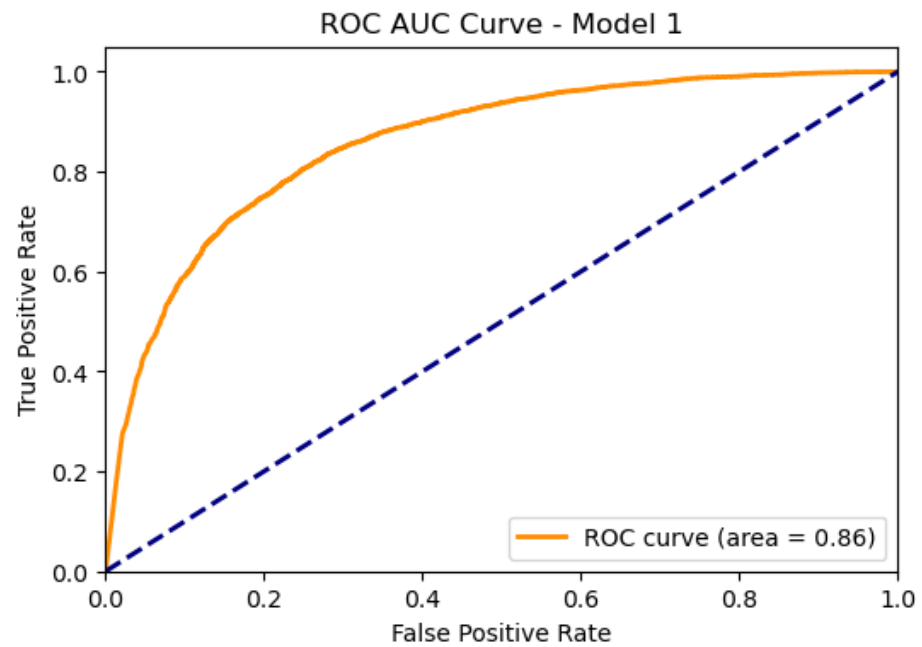
# Predict probabilities for both models
y_score_1 = model_1.predict(x_test)
y_score_2 = model_2.predict(x_test)

# Threshold for binary classification
y_pred_1 = (y_score_1 > 0.5)
y_pred_2 = (y_score_2 > 0.5)

```

```
625/625 [=====] - 2s 2ms/step  
625/625 [=====] - 2s 3ms/step
```







```
In [ ]: # Plot confusion matrix and ROC AUC curve for Model 1
plot_confusion_matrix(y_test, y_pred_1, title="Confusion Matrix - Model 1")
plot_roc_auc(y_test, y_score_1, title="ROC AUC Curve - Model 1")
```



In [28]:

```
# Plot confusion matrix and ROC AUC curve for Model 2  
plot_confusion_matrix(y_test, y_pred_2, title="Confusion Matrix - Model 2")  
plot_roc_auc(y_test, y_score_2, title="ROC AUC Curve - Model 2")
```





Certainly! Let's analyze the classification reports for Model 1 and Model 2:

Model 1: Precision for Class 0 (0.68): This means that when Model 1 predicts a sample as Class 0, it is correct 68% of the time. Recall for Class 0 (0.65): Model 1 correctly identifies 65% of all actual Class 0 samples. F1-Score for Class 0 (0.66): The F1-score is the harmonic mean of precision and recall. It gives a balanced measure of a model's performance. An F1-score of 0.66 for Class 0 indicates a reasonable balance between precision and recall. Precision for Class 1 (0.90): Model 1's precision for Class 1 is quite high, indicating that when it predicts Class 1, it is correct 90% of the time. Recall for Class 1 (0.91): Model 1 correctly identifies 91% of all actual Class 1 samples. F1-Score for Class 1 (0.90): The F1-score for Class 1 is high, indicating a strong balance between precision and recall for Class 1. Accuracy (0.85): Model 1's overall accuracy is 85%, meaning it correctly predicts 85% of all samples. Macro Avg F1-Score (0.78): The macro-average F1-score is the average of the F1-scores for both classes. It gives equal weight to both classes. In this case, it's 0.78. Weighted Avg F1-Score (0.85): The weighted average F1-score takes into account class imbalance. Since Class 1 has more samples, it has a greater influence on the weighted average. The weighted F1-score is 0.85.

Model 2: Precision for Class 0 (0.62): Model 2's precision for Class 0 is lower compared to Model 1, indicating that it is less accurate when predicting Class 0. Recall for Class 0 (0.61): Model 2 correctly identifies 61% of all actual Class 0 samples. F1-Score for Class 0 (0.61): The F1-score for Class 0 is relatively low, indicating a trade-off between precision and recall. Precision for Class 1 (0.88): Model 2's precision for Class 1 is high, indicating that it is accurate when predicting Class 1. Recall for Class 1 (0.89): Model 2 correctly identifies 89% of all actual Class 1 samples. F1-Score for Class 1 (0.89): The F1-score for Class 1 is high, indicating a good balance between precision and recall for Class 1. Accuracy (0.82): Model 2's overall accuracy is 82%, which is slightly lower than Model 1. Macro Avg F1-Score (0.75): The macro-average F1-score for Model 2 is 0.75, indicating that it performs slightly worse in terms of overall balance between precision and recall compared to Model 1. Weighted Avg F1-Score (0.82): The weighted average F1-score is 0.82, which considers class imbalance and is higher than the macro-average F1-score. Analysis:

Model 1 outperforms Model 2 in terms of precision, recall, and F1-scores for both classes (0 and 1). Model 2 has lower overall accuracy compared to Model 1. Model 2 has a slight imbalance in class performance, with Class 1 being predicted more accurately than Class 0. Model 1 is generally a better-performing model based on these classification metrics and should be preferred if you prioritize balanced performance between the two classes. However, further analysis, such as feature importance or

domain-specific considerations, may be necessary to make a final decision on model selection.

In [ ]: